

Lecture 2

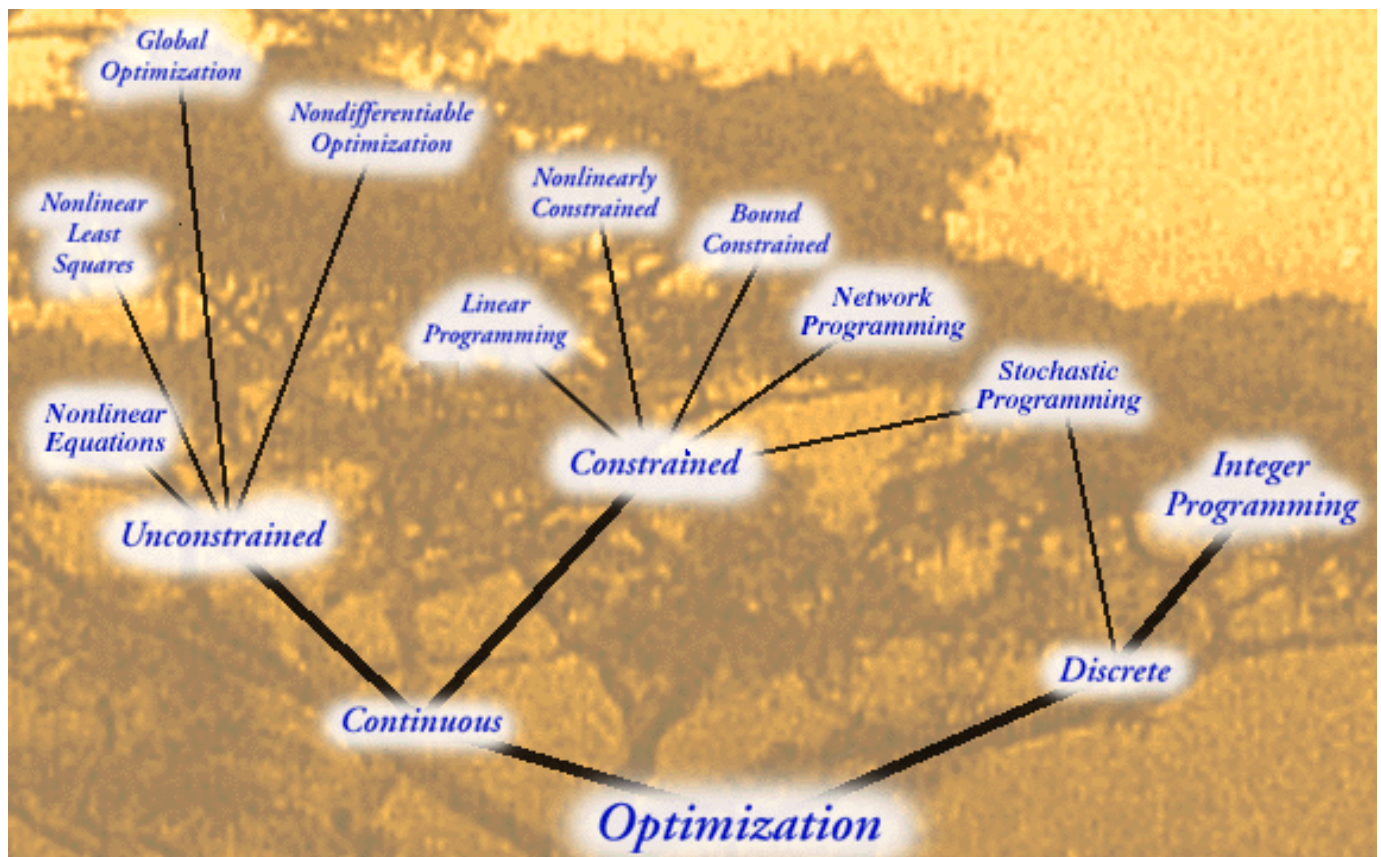
C25 Optimization

Hilary 2013

A. Zisserman

- Discrete optimization
- Dynamic Programming
- Applications
- Generalizations ...

The Optimization Tree



Discrete Optimization

- Also called combinatorial optimization
- In continuous optimization each component, x_i , of the vector \mathbf{x} can vary continuously
- In discrete optimization each component, or variable, x_i can only take a finite number of possible states
- The advantage is that for special cases the global optimum can be found for cost functions which are not convex when the variables are continuous
- And these special cases occur often in real applications
- And the global optimum can be found efficiently

—————
continuous

—|—|—|—|—|—
discrete

Dynamic programming


- Applies to problems where the cost function can be:
 - decomposed into a sequence (ordering) of stages, and
 - each stage depends on only a fixed number of previous stages
- The cost function need not be convex (if variables continuous)
- The name “dynamic” is historical
- Also called the “Viterbi” algorithm

Consider a cost function $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$ of the form

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi_i(x_{i-1}, x_i)$$

where x_i can take one of h values

e.g. $h=5, n=6$




Complexity of minimization:

- exhaustive search $O(h^n)$
- dynamic programming $O(nh^2)$

Example 1

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$

$$f(\mathbf{x}) = \sum_{i=1}^n \underbrace{(x_i - d_i)^2}_{\text{closeness to measurements}} + \sum_{i=2}^n \underbrace{\lambda^2 (x_i - x_{i-1})^2}_{\text{smoothness}}$$



Motivation: complexity of stereo correspondence

Objective: compute horizontal displacement for matches between left and right images




x_i is spatial shift of i 'th pixel $\rightarrow h = 40$

\mathbf{x} is all pixels in row $\rightarrow n = 256$

Complexity $O(40^{256})$ vs $O(256 \times 40^2)$

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$



Key idea: the optimization can be broken down into n sub-optimizations


Step 1: For each value of x_2 determine the best value of x_1

- Compute

$$\begin{aligned} S_2(x_2) &= \min_{x_1} \{m_2(x_2) + m_1(x_1) + \phi(x_1, x_2)\} \\ &= m_2(x_2) + \min_{x_1} \{m_1(x_1) + \phi(x_1, x_2)\} \end{aligned}$$

- Record the value of x_1 for which $S_2(x_2)$ is a minimum

To compute this minimum for all x_2 involves $O(h^2)$ operations



Step 2: For each value of x_3 determine the best value of x_2 **and** x_1

- Compute

$$S_3(x_3) = m_3(x_3) + \min_{x_2} \{S_2(x_2) + \phi(x_2, x_3)\}$$

- Record the value of x_2 for which $S_3(x_3)$ is a minimum

Again, to compute this minimum for all x_3 involves $O(h^2)$ operations
 Note $S_k(x_k)$ encodes the lowest cost partial sum for all nodes up to k which have the value x_k at node k , i.e.

$$S_k(x_k) = \min_{x_1, x_2, \dots, x_{k-1}} \sum_{i=1}^k m_i(x_i) + \sum_{i=2}^k \phi(x_{i-1}, x_i)$$

Viterbi Algorithm

- Initialize $S_1(x_1) = m_1(x_1)$

- For $k = 2 : n$

$$S_k(x_k) = m_k(x_k) + \min_{x_{k-1}} \{S_{k-1}(x_{k-1}) + \phi(x_{k-1}, x_k)\}$$

$$b_k(x_k) = \arg \min_{x_{k-1}} \{S_{k-1}(x_{k-1}) + \phi(x_{k-1}, x_k)\}$$

- Terminate

$$x_n^* = \arg \min_{x_n} S_n(x_n)$$

- Backtrack

$$x_{i-1} = b_i(x_i)$$



Complexity $O(nh^2)$

Example 2

$$f(\mathbf{x}) = \sum_{i=1}^n (x_i - d_i)^2 + \sum_{i=2}^n g_{\alpha, \lambda}(x_i - x_{i-1})$$

where

$$g_{\alpha, \lambda}(\Delta) = \min(\lambda^2 \Delta^2, \alpha) = \begin{cases} \lambda^2 \Delta^2 & \text{if } |\Delta| < \sqrt{\alpha}/\lambda \\ \alpha & \text{otherwise.} \end{cases}$$



Note, $f(\mathbf{x})$ is not convex

Note

This type of cost function often arises in MAP estimation

$$\mathbf{x}^* = \arg \max_{\mathbf{x}} p(\mathbf{x}|\mathbf{y})$$

measurements

$$= \arg \max_{\mathbf{x}} p(\mathbf{y}|\mathbf{x})p(\mathbf{x})$$

Bayes' rule

$$\sim \prod_i^n e^{-\frac{(x_i - y_i)^2}{2\sigma^2}} e^{-\beta^2(x_i - x_{i-1})^2}$$

e.g. for Gaussian measurement errors, and first order smoothness

Use negative log to obtain a cost function of the form

$$f(\mathbf{x}) = \sum_{i=1}^n \underbrace{(x_i - y_i)^2}_{\text{from likelihood}} + \sum_{i=2}^n \underbrace{\lambda^2(x_i - x_{i-1})^2}_{\text{from prior}}$$

Where can DP be applied?

Dynamic programming can be applied when there is a linear ordering on the cost function (so that partial minimizations can be computed).

Example Applications:

1. Text processing: String edit distance
2. Speech recognition: Dynamic time warping
3. Computer vision: Stereo correspondence
4. Image manipulation: Image re-targeting
5. Bioinformatics: Gene alignment
6. Hidden Markov model (HMM)

Application I: string edit distance

The **edit distance** of two strings, s_1 and s_2 , is the minimum number of single character mutations required to change s_1 into s_2 , where a mutation is one of:

1. substitute a letter (kat \rightarrow cat) cost = 1
2. insert a letter (ct \rightarrow cat) cost = 1
3. delete a letter (caat \rightarrow cat) cost = 1

Example: $d(\text{opimizeon}, \text{optimization})$

```
op imizeon
|| |||||
optimization
|||||
cciccccccc
```

'c' = copy, cost = 0

$d(s_1, s_2) = 2$

Complexity

- for two strings of length m and n , exhaustive search has complexity $O(3^{m+n})$
- dynamic programming reduces this to $O(mn)$

Using string edit distance for spelling correction

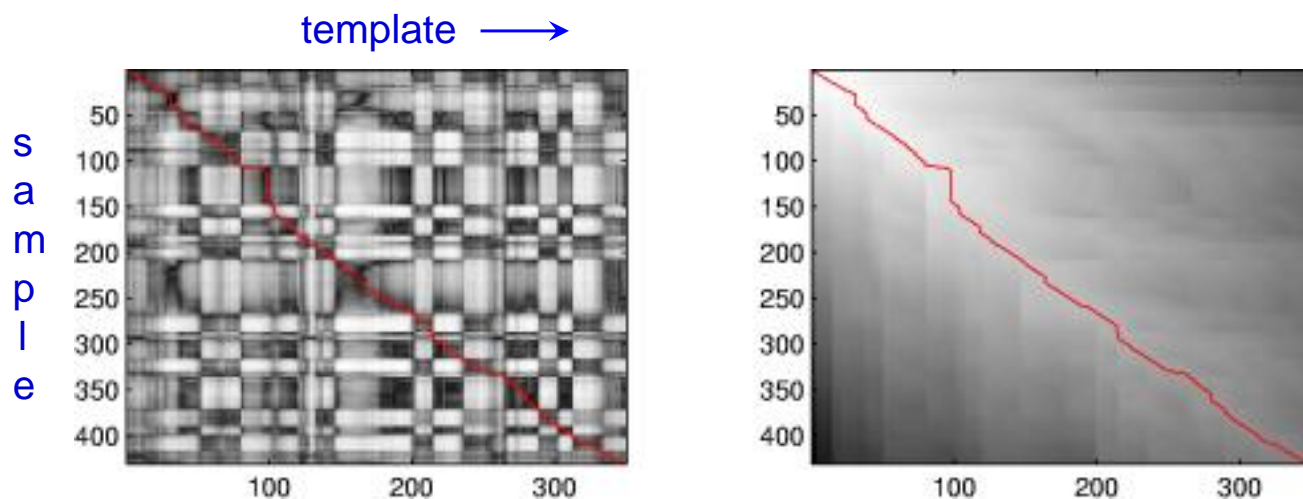
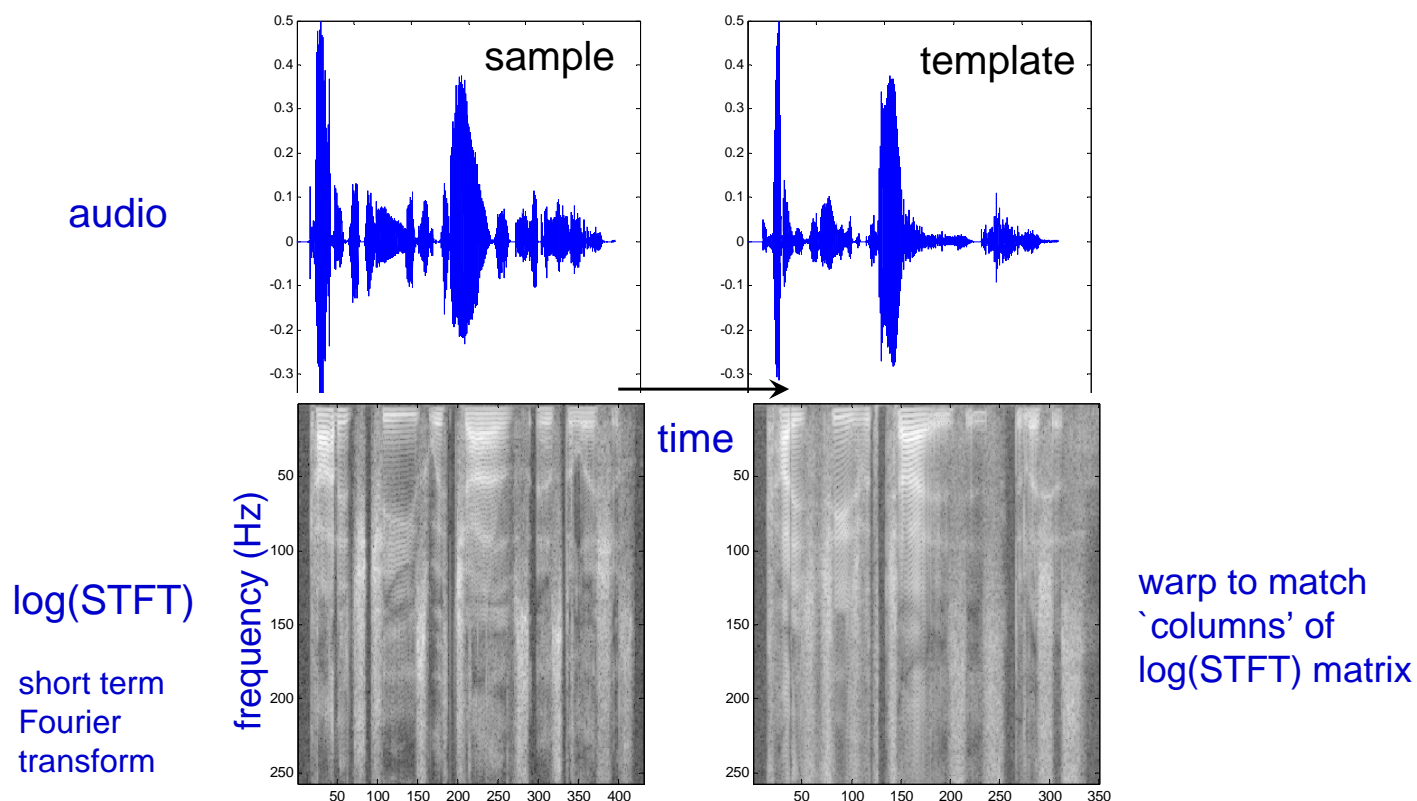
1. Check if word w is in the dictionary D
2. If it is not, then find the word x in D that minimizes $d(w, x)$
3. Suggest x as the corrected spelling for w

Note: step 2 appears to require computing the edit distance to all words in D , but this is not required at run time because edit distance is a **metric**, and this allows efficient search.

Mispelling	Word	ED	Detail
committment	commitment	1	committment \rightarrow commitment
tommorrow	tomorrow	1	tommorrow \rightarrow tomorrow
saftey	safety	2	saftey \rightarrow safty \rightarrow safety

Application II: Dynamic Time Warp (DTW)

Objective: temporal alignment of a sample and template speech pattern



x_i is time shift of i th column

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$

quality of match \nearrow cost of allowed moves \nearrow

$\rightarrow (1, 0)$
 $\downarrow (0, 1)$
 $\searrow (1, 1)$

Application III: stereo correspondence

Objective: compute horizontal displacement for matches between left and right images

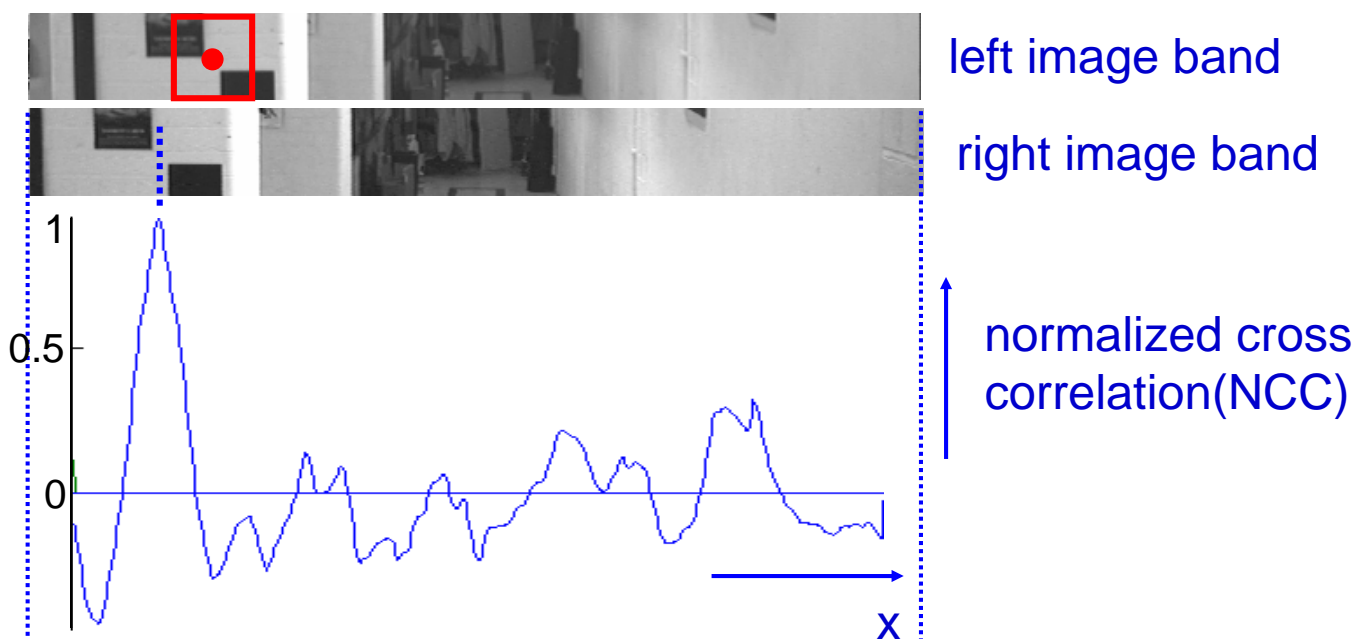


x_i is spatial shift of i th pixel

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$

quality of match \nearrow uniqueness, smoothness \nearrow


$$m(x) = \alpha(1 - \text{NCC})^2$$



NCC of square image regions at offset (disparity) x

-


left image



Real-time application – Background substitution



Left view



Right view

Input

Results



input left view




Background substitution 1



Background substitution 2

Application IV: image re-targeting

- Remove image “seams” for imperceptible aspect ratio change





scale




seam
removal



Finding the optimal seam – s

$E(I)$



$$E(I) = |\partial I / \partial x| + |\partial I / \partial y| \rightarrow s^* = \arg \min_s E(s)$$


Dynamic Programming on graphs

- **Graph** (V, E)
- **Vertices** v_i for $i = 1, \dots, n$
- **Edges** e_{ij} connect v_i to other vertices v_j

$$f(\mathbf{x}) = \sum_{v_i \in V} m_i(v_i) + \sum_{e_{ij} \in E} \phi(v_i, v_j)$$


e.g.

- 3 vertices, each can take one of h values
- 3 edges




Dynamic Programming on graphs


So far have considered **chains**




Can dynamic programming be applied to these configurations?



Fully connected



Tree structure



Star structure